

A Cybernetics-Inspired Adaptive Trading System: Mathematical Methods and Implementation

July 20, 2025

Contents

1	Introduction	1
2	Data and Notation	1
3	Feature Engineering	1
3.1	Returns	1
3.2	Rolling Mean and Standard Deviation	1
3.3	Rolling Entropy	2
4	Information-Theoretic Feature Selection: Mutual Information	2
5	Signal Denoising: Wiener Filtering	2
6	Adaptive Weights: Online Linear Learning	3
7	Feedback-Based Control Law	3
8	Backtesting Engine	3
8.1	Simulation Protocol	3
8.2	Performance Metrics	4
9	Implementation Notes and Bias Avoidance	4
10	Summary	4

1 Introduction

This document details the theoretical and practical implementation of a cybernetic trading and forecasting framework, inspired by Norbert Wiener’s *Cybernetics: Or Control and Communication in the Animal and the Machine*. The approach integrates core concepts from information theory, control theory, adaptive systems, and quantitative finance—most notably, rolling entropy, mutual information, Wiener filtering, feedback-based adaptive control, and a robust backtesting protocol, all implemented with awareness of look-ahead bias.

2 Data and Notation

Let P_t denote the price (e.g., close) of an asset at time t . Assume a time series P_1, P_2, \dots, P_T sampled at regular intervals (e.g., daily). Let X_t denote a vector of features or indicators constructed using the series up to time t .

3 Feature Engineering

Features are designed to capture both statistical properties and market regimes. All series and operations are computed using **only past and present values** (never future), to avoid forward-looking bias.

3.1 Returns

$$r_t = \frac{P_t - P_{t-1}}{P_{t-1}}$$

3.2 Rolling Mean and Standard Deviation

Define, for window length w :

$$\text{Mean}_t = \frac{1}{w} \sum_{i=t-w+1}^t P_i$$
$$\text{Std}_t = \sqrt{\frac{1}{w} \sum_{i=t-w+1}^t (P_i - \text{Mean}_t)^2}$$

3.3 Rolling Entropy

Given a vector $V = (P_{t-w+1}, \dots, P_t)$ of length w , construct a histogram with b bins, yielding empirical probabilities p_j .

$$\text{Entropy}_t = - \sum_{j=1}^b p_j \log_2 p_j$$

where p_j is the normalized frequency of samples in bin j , and we set $0 \log_2 0 = 0$ by convention.

4 Information-Theoretic Feature Selection: Mutual Information

To quantify how much information a candidate feature carries about the future return, we use mutual information. Given feature matrix X (shape $n \times m$) and future asset return y (length n):

$$\text{MI}(X_j; y) = \sum_{x,y} p(x,y) \log \frac{p(x,y)}{p(x)p(y)}$$

where X_j is the j -th feature. Features with the highest MI are selected as primary predictors for the next window.

5 Signal Denoising: Wiener Filtering

For a time series z_t with additive noise, the Wiener filter estimates the underlying "clean" signal using past and present values.¹

$$\hat{z}_t = \mathcal{W}(z_{t-k}, \dots, z_t; \sigma_n^2)$$

where \mathcal{W} denotes the Wiener filter (e.g., implemented as a moving window, minimizing mean squared error), and σ_n^2 is the estimated noise variance.

6 Adaptive Weights: Online Linear Learning

Let each signal X_t (for t in the look-back window) be an m -dimensional feature vector. We aim to estimate a weight vector \mathbf{w} for predicting the future target y_t (such as next step return):

$$\hat{y}_t = \mathbf{w}^\top X_t$$

Weights are updated using **online gradient descent** (stochastic gradient):

$$\mathbf{w}_{k+1} = \mathbf{w}_k + \eta(y_k - \hat{y}_k)X_k$$

where $\eta > 0$ is a learning rate, y_k is the target (historical observed return), and \hat{y}_k is the predicted value using current weights.

7 Feedback-Based Control Law

Inspired by regulatory mechanisms in cybernetics and control engineering, the trading action is based on the feedback between the desired and predicted signal. For proportional control (P-controller):

$$\text{TradeSignal}_t = K_p(T - S_t)$$

where T is a target threshold (typically zero), S_t is the predicted signal (e.g., predicted return), and K_p is proportional gain (often set to 1). The trade decision at time t is ± 1 (long/short) based on the sign of TradeSignal_t .

8 Backtesting Engine

8.1 Simulation Protocol

At each step t :

¹See Wiener, Norbert, *Cybernetics*, Ch. III.

Step 1: Use only data up to (and including) time $t - 1$ to compute features, filter, adapt weights, select predictors, and generate trade signal.

Step 2: Observe price P_t and determine P&L from holding previous position over $[t - 1, t]$.

Step 3: Execute trade at P_t , taking into account transaction costs:

$$\text{Cost}_t = \gamma |\text{Position}_t - \text{Position}_{t-1}| \cdot P_t$$

where γ is the proportional trading fee.

Step 4: Update equity by

$$\text{Equity}_t = \text{Cash}_{t-1} - \text{Cost}_t + \text{Position}_{t-1}(P_t - P_{t-1})$$

and $\text{Cash}_t = \text{Equity}_t - \text{Position}_t \cdot P_t$.

8.2 Performance Metrics

Backtest results are evaluated via:

- **Total Return:** $\frac{\text{Equity}_T}{\text{Equity}_0} - 1$
- **Sharpe Ratio:** $\frac{\mu_R}{\sigma_R} \sqrt{N}$, where R is the periodic return
- **Maximum Drawdown:** $\min_t \left(\frac{\text{Equity}_t}{\max_{s \leq t} \text{Equity}_s} - 1 \right)$

9 Implementation Notes and Bias Avoidance

- **No Data Leakage:** All computations for trading decisions at time t use features and targets computed from $[t - \text{lookback}, t - 1]$ only—i.e., strictly up to, not including, time t 's outcome.
- **Synchronizing Indices:** Feature vectors and price series are aligned and dropped where missing data could introduce misaligned targets.
- **Online Adaptation:** Weights are updated per time window, based on rolling, in-sample past data only.
- **Transaction Costs:** Simulated on all position changes, to reflect realistic implementation.

10 Summary

This cybernetics-inspired trading system provides an integrated workflow for interpretable, adaptive algorithmic trading. Methods include:

- Rolling and information-theoretic feature engineering,
- Predictive signal denoising via Wiener filters,

- Adaptive online model learning via feedback control,
- Bias-avoiding and reproducible backtesting.

Such systems reflect Wiener's vision of robust, feedback-driven control and inference under uncertainty, applied here to modern algorithmic financial markets.